

GPT-4.1 Mini: robustness under prompt paraphrase  
(spread across prompts = 35.9% of baseline)

